




IBM Informix

Improving IDS Administration Zone

Administration Free Zone

- **Low Total Cost of Ownership and Deployment**
- **Easy Scalable Administration**
- **Open Admin API for customized administration**
- **Autonomic architecture**
- **Industrial strength, highly reliable - Install it, Set it up, and Forget about it.**



Overview

- **SQL Based Administration Commands**
- **Improved Sysmaster Database**
- **SQL History Tracing**
- **Built in Database Scheduler**
- **Quick overview of IDSAdmin**



Laying the Foundation for Improvements


- **A new database (sysadmin) for**
 - Administrative functions
 - Alert System
 - Data collection
- **Different methods for collecting information**
- **Additional information required to present in a way DBAs can understand**



Improve Information for sysmaster and onstat

- **Improve the display of information in the following areas:**

- Improved Thread Wait Information (onstat -g ath)
- Checkpoint Information (onstat -g ckp)
- User/Server Environment Information (onstat -g env)
- Online & onbar log files (onstat -m)
- Storage system IO counts and times (onstat -g iof)
- MGM Information (onstat -g mgm)
- SQL Generic Cache Profiles (onstat -g ssc)
- SQL Statement Performance History (onstat -g his)



Make the information clear for the DBA to understand

New Sysmaster Tables

| TABLE NAME | DESCRIPTION |
|--------------------------|---|
| syscheckpoint | The information about the checkpoint and associated statistics |
| sys tcblst | Modified the existing table to add wait stats. |
| sysenvses | View Informix's session environment variables |
| sysenv | View the servers environment variables |
| sysonline log | View the online.log for the server |
| sys scblst | Improvement to view the memory used by session |
| sysnetworkio | View the network I/O generated by database session |
| sysdual | Oracle compatibility feature |
| sys sqlcacheprof | Displays the profile information about each SQL cache |
| sys sqltrace | The sql statements which have been recently executed on the system |
| sys sqltrace_itr | The list of iterators for the SQL statement. |
| sys sqltrace_info | General information about the SQL tracing |
| sysnetglobal | Global Network Information |
| sysnetclienttype | Network information based on client type |
| sysbaract_log | The OnBar Activity Log file |
| sysrstcb | Improvement to view I/O and lock wait information |

Onstat -g ath

- Remove/Reduce the number of thread status of “sleeping forever”
- Give the DBA a clear picture of what is happening

Threads:

| tid | tcb | rstcb | prty | status | vp-class | name |
|-----|-----------|-----------|------|-------------------|----------|--------------|
| 2 | 10bbf36a8 | 0 | 2 | sleeping forever | 3lio | lio vp 0 |
| 3 | 10bc12218 | 0 | 2 | sleeping forever | 4pio | pio vp 0 |
| 4 | 10bc31218 | 0 | 2 | sleeping forever | 5aio | aio vp 0 |
| 5 | 10bc50218 | 0 | 2 | sleeping forever | 6msc | msc vp 0 |
| 6 | 10bc7f218 | 0 | 2 | sleeping forever | 7aio | aio vp 1 |
| 7 | 10bc9e540 | 10b231028 | 4 | sleeping secs: 1 | 1cpu | main_loop() |
| 8 | 10bc12548 | 0 | 2 | running | 1cpu | tlitcppoll |
| 9 | 10bc317f0 | 0 | 3 | sleeping forever | 1cpu | tlitcplst |
| 10 | 10bc50438 | 10b231780 | 2 | sleeping forever | 1cpu | flush_sub(0) |
| 11 | 10bc7f740 | 0 | 2 | sleeping forever | 8aio | aio vp 2 |
| 12 | 10bc7fa00 | 0 | 2 | sleeping forever | 9aio | aio vp 3 |
| 13 | 10bd56218 | 0 | 2 | sleeping forever | 10aio | aio vp 4 |
| 14 | 10bd75218 | 0 | 2 | sleeping forever | 11aio | aio vp 5 |
| 15 | 10bd94548 | 10b231ed8 | 3 | sleeping forever | 1cpu | aslogflush |
| 16 | 10bc7fd00 | 10b232630 | 1 | sleeping secs: 26 | 1cpu | btscanner 0 |
| 32 | 10c738ad8 | 10b233c38 | 4 | sleeping secs: 1 | 1cpu | onmode_mon |
| 50 | 10c0db710 | 10b232d88 | 2 | cond wait netnorm | 1cpu | sqlxec |

Improved onstat -g ath

Threads:

| tid | tcb | rstcb | prty | status | vp-class | name |
|-----|-----------|-----------|------|-------------------|----------|--------------|
| 2 | 10bbf36a8 | 0 | 2 | IO Idle | 3lio | lio vp 0 |
| 3 | 10bc12218 | 0 | 2 | IO Idle | 4pio | pio vp 0 |
| 4 | 10bc31218 | 0 | 2 | running | 5aio | aio vp 0 |
| 5 | 10bc50218 | 0 | 2 | IO Idle | 6msc | msc vp 0 |
| 6 | 10bc7f218 | 0 | 2 | running | 7aio | aio vp 1 |
| 7 | 10bc9e540 | 10b231028 | 4 | sleeping secs: 1 | 1cpu | main_loop() |
| 8 | 10bc12548 | 0 | 2 | running | 1cpu | tlitcpoll |
| 9 | 10bc317f0 | 0 | 3 | sleeping forever | 1cpu | tlitcplst |
| 10 | 10bc50438 | 10b231780 | 2 | IO Wait | 1cpu | flush_sub(0) |
| 11 | 10bc7f740 | 0 | 2 | IO Idle | 8aio | aio vp 2 |
| 12 | 10bc7fa00 | 0 | 2 | IO Idle | 9aio | aio vp 3 |
| 13 | 10bd56218 | 0 | 2 | IO Idle | 10aio | aio vp 4 |
| 14 | 10bd75218 | 0 | 2 | IO Idle | 11aio | aio vp 5 |
| 15 | 10bd94548 | 10b231ed8 | 3 | sleeping forever | 1cpu | aslogflush |
| 16 | 10bc7fd00 | 10b232630 | 1 | sleeping secs: 34 | 1cpu | btscanner 0 |
| 32 | 10c738ad8 | 10b233c38 | 4 | sleeping secs: 1 | 1cpu | onmode_mon |
| 50 | 10c0db710 | 10b232d88 | 2 | IO Wait | 1cpu | sqlxec |

SQL Admin Commands

- **A set of User Defined Routines (UDRs) to administer the Informix database server.**
- **The major categories of administration include:**
 - Space Management
 - Configuration Management
 - Routine task maintenance
 - System Validation (oncheck functionality)
- **Feature Benefits**
 - SQL Based Administration
 - Remote Administration
 - Tracking of command execution and results in a system table



Admin Commands – Two New UDRs

- Two UDRs called **task** & **admin** are part of the sysadmin database
- They perform exactly the same, only the return code is different
 - *task()* UDR returns a character string describing the return status

```
EXECUTE FUNCTION task('create dbspace', 'dbspace2', '/CHUNKS/dbspace2');  
(expression)  created dbspace number 2 named dbspace2
```

- *admin()* UDR returns an integer which is a link to the *command_history* table

```
EXECUTE FUNCTION admin('create dbspace', 'dbspace2', '/CHUNKS/dbspace2');  
(expression)  107
```

- Both UDRS log all executions into a table called *command_history* in the *sysadmin* database

SQL Admin Command - Tracking

- **SQL Admin API logs all executions into a table called *command_history* in the *sysadmin* database**

| Column | Type | Description |
|----------------|-------------------------------|--|
| cmd_number | serial | |
| cmd_exec_time | datetime year to second | Time the command was started |
| cmd_user | varchar | User executing the command |
| cmd_hostname | varchar | Host the command was executed from |
| cmd_executed | varchar | The command executed |
| cmd_ret_status | Integer | Return code |
| cmd_ret_msg | lvarchar | Return message |

SQL Admin Commands - Parameters

- **Environment Variable Expansion**

- A pathname may start with an environment variable.
- The environment variable may only exist in the server's environment

- **Unit Extensions**

- All offsets and sizes can be provided with unit extensions
- The extensions are case insensitive
- Default is KB
- PB, TB, GB, MB, KB, B

SQL Admin Commands – EXAMPLE

```
EXECUTE FUNCTION admin('create dbspace','dbspace2','$INFORMIXDIR/SPACE/dbspace2','20MB');  
(expression)          108
```

```
SELECT * FROM command_history WHERE cmd_number IN (108)
```

| | |
|----------------|---|
| cmd_number | 108 |
| cmd_exec_time | 2005-11-17 16:26:15 |
| cmd_user | informix |
| cmd_hostname | olympia.beaverton.ibm.com |
| cmd_executed | create dbspace |
| cmd_ret_status | 0 |
| cmd_ret_msg | created dbspace number 2 named dbspace2 |

Admin API Commands

- ADD BUFFERPOOL
- ADD CHUNK
- ADD LOG
- ADD MEMORY
- ADD MIRROR
- ALTER CHUNK OFFLINE
- ALTER CHUNK ONLINE
- ALTER LOGMODE
- ALTER PLOG
- ARCHIVE FAKE
- CHECK DATA
- CHECK EXTENTS
- CHECK PARTITION
- CHECKPOINT
- CLEAN SBSPACE
- CREATE BLOBSpace
- CREATE CHUNK
- CREATE DBSPACE
- CREATE SBSPACE
- CREATE TEMPDBSPACE
- DROP BLOBSpace
- DROP CHUNK
- DROP DBSPACE
- DROP LOG
- DROP SBSPACE
- DROP TEMPDBSPACE
- ONMODE
- PRINT ERROR
- PRINT PARTITION
- QUIESCENT
- RENAME SPACE
- SET CHUNK OFFLINE
- SET CHUNK ONLINE
- SET DATASKIP ON
- SET DATASKIP OFF
- SET SBSPACE ACESSTIME ON
- SET SBSPACE ACESSTIME OFF
- SET SBSPACE AVG_LO_SIZE
- SET SBSPACE LOGGING ON
- SET SBSPACE LOGGING OFF
- SET SQL TRACING
- SET SQL TRACING OFF
- SET SQL TRACING ON
- SET SQL TRACING RESIZE
- SET SQL USER TRACING
- SET SQL USER TRACING CLEAR
- SET SQL USER TRACING OFF
- SHUTDOWN
- START MIRRORING space
- STOP MIRRORING

See IBM Informix Guide to SQL: Syntax, Chapter 6

DBA Constant Struggle

Identifying Performance Bottlenecks in SQL Statements

- **Current methods for examination**
 - Set explain
 - Command line utilities
 - sysmaster
- **Lack of simplicity in the process**
- **Hard to build a repeatable process**



Questions DBAs like to ask?

- **How long did a SQL statements take?**
- **How many resources of each category did a statement take?**
 - Disk I/O
 - Memory
 - CPU
- **How long and how many times did we wait on each resource?**
 - Locks
 - Disk I/O



Solutions – SQL Query Drill Down Feature

- **Collect detailed sql performance information.**
- **Display using onstat or sysmaster.**
- **Dynamically configurable**
- **By default disabled**
- **Global and User Tracing modes**

| Name | Description |
|--------------------------|--|
| User ID | User id |
| Session ID | Database Session ID |
| Database Name | Current Database Name |
| Statement Type | Type of SQL statement being executed |
| Statement Execution time | Duration of the SQL statement |
| Time of Execution | Date & Time this statement completed |
| Statement Text | SQL statement text or the procedure stack trace with statement type |
| RSAM statistics | <ul style="list-style-type: none">• Buffer Reads & Writes• Page reads & Writes• Memory Sorts, disk Sorts,• Lock requests, waits• Logical Log Records |
| SQL statistics | <ul style="list-style-type: none">• Estimated # of rows• Estimated Cost• # of rows returned• Statement Type• Database Isolation Level |

Controlling SQL Query Drill Down

- **ONCONFIG variable SQLTRACE**

- Level =[off,low,med,high]
- Ntraces=[number of traces]
- Size=[size of each trace buffer in KB]

```
SQLTRACE level=low,ntraces=2000,size=1,mode=global
```

- **Turn off SQL Tracing for session id 147**

```
execute function task("SET SQL USER TRACING OFF",147);
```

Controlling SQL Query Drill Down

- **Dynamically enable or modify SQL Tracing**

- Trace 2000 SQL statements
- Trace 1024 bytes of data for each SQL statement

```
execute function task("SET SQL TRACING ON",2000,1);
```

- **Turn off SQL Tracing**

```
execute function task("SET SQL TRACING OFF");
```

SQL Query Drill Down - onstat

Database: sysadmin

Statement text:

```
SELECT MAX(run_task_seq) FROM ph_run A, ph_task B WHERE A.run_task_id = ?
AND A.run_task_id = B.tk_id AND A.run_time + B.tk_delete < CURRENT
```

Iterator/Explain

=====

| ID | Left | Right | Est Cost | Est Rows | Num Rows | Type |
|----|------|-------|----------|----------|----------|-------------|
| 3 | 0 | 0 | 1 | 1 | 1 | Index Scan |
| 4 | 0 | 0 | 19 | 545 | 1 | Index Scan |
| 2 | 3 | 4 | 20 | 5 | 1 | Nested Join |
| 1 | 2 | 0 | 1 | 1 | 1 | Group |

Statement information:

| Sess_id | User_id | Stmt Type | Finish Time | Run Time |
|---------|---------|-----------|-------------|----------|
| 21 | 0 | SELECT | 10:51:11 | 0.0023 |

Statement Statistics:

| Page | Buffer | Read | Buffer | Page | Buffer | Write |
|------|--------|---------|----------|-------|--------|---------|
| Read | Read | % Cache | IDX Read | Write | Write | % Cache |
| 0 | 77 | 100.00 | 0 | 0 | 0 | 0.00 |


| Lock | Lock | LK Wait | Log | Num | Disk | Memory |
|----------|-------|----------|---------|-------|-------|--------|
| Requests | Waits | Time (S) | Space | Sorts | Sorts | Sorts |
| 0 | 0 | 0.0000 | 0.000 B | 0 | 0 | 0 |

| Total | Total | Avg | Max | Avg | I/O Wait | Avg Rows |
|------------|----------|----------|----------|----------|----------|----------|
| Executions | Time (S) | Time (S) | Time (S) | IO Wait | Time (S) | Per Sec |
| 220 | 78.8463 | 0.3584 | 1.9557 | 0.000000 | 0.000000 | 439.9908 |

| Estimated | Estimated | Actual | SQL | ISAM | Isolation | SQL |
|-----------|-----------|--------|-------|-------|-----------|--------|
| Cost | Rows | Rows | Error | Error | Level | Memory |
| 20 | 1 | 1 | 0 | 0 | DR | 41552 |

Built-in Database Scheduler

- **Ability to schedule SQL, Stored procedures or a UDR**
 - Timed backup, profile information.
- **Scheduler managed entities:**
 - Tasks
 - Sensor
 - Startup Task
 - Startup Sensor
- **These entities are driven by the data inside a table called *ph_task*. Each row is a task.**



You can schedule your Tasks on the following base:

- **Schedule to only run once at server startup time**
- **Schedule at a specific start time**
- **Schedule in a specified frequency**
- **Schedule a combination of start time and frequency**

Scheduler Specialized Task

■ Task

- A means to execute a specific job at a specific time or interval
- No intention to evaluate the returned data
- The *task* executes by invoking
 - A single or compound SQL statement
 - Stored procedure,
 - C User Defined Routine
 - Java User Defined Routine



Sensor - Specialized Task

- **A specialized task geared at collecting and saving data without the DBA doing much work.**
- **A means to get information about a managed element**

Why use Tasks and Sensors?

- **Build a specialized task**
 - A way of ensuring routine jobs get completed
 - Periodically check, analyze data server is operating efficiently
- **Sensors provide a simple way of collecting information**
 - Easy to add a new sensor
 - Provide a portable way of collecting information without using the operating system.



Examples of Creating a Sensor

- **Create a sensor to track the number of sessions on the database server every 5 minutes.**
- **Only keep the data for 1 day**

```
INSERT INTO ph_task (tk_name,tk_type,tk_group,tk_description,tk_result_table,  
tk_create,tk_execute,tk_start_time,tk_stop_time,tk_frequency,tk_delete)  
VALUES  
("mon_user_count","SENSOR","SERVER","Count the number of user count ",  
"mon_user_count","create table mon_user_count (ID integer, session_count integer)",  
"insert into mon_user_count select $DATA_SEQ_ID, count(*) from  
    sysmaster:syssessions",  
NULL,NULL,  
INTERVAL ( 5 ) MINUTE TO MINUTE,  
INTERVAL ( 1 ) DAY TO DAY);
```

The PH_TASK table

```
select * from PH_TASK where tk_name="mon_user_count"
```


| | |
|-----------------|--|
| tk_id | 19 |
| tk_name | mon_user_count |
| tk_description | Count the number of user count |
| tk_type | SENSOR |
| tk_sequence | 243 |
| tk_result_table | mon_user_count |
| tk_create | create table mon_user_count (ID integer, session_count integer) |
| tk_dbs | sysadmin |
| tk_execute | insert into mon_user_count select \$DATA_SEQ_ID, count(*) from sysmaster:syssessions |

| | |
|--------------------|---------------------|
| tk_delete | 30 00:00:00 |
| tk_start_time | |
| tk_stop_time | |
| tk_frequency | 0 00:05:00 |
| tk_next_execution | 2007-07-15 07:19:21 |
| tk_total_execution | 2 |
| tk_total_time | 2.3839 |
| tk_monday | t |
| tk_tuesday | t |
| tk_wednesday | t |
| tk_thursday | t |
| tk_friday | t |
| tk_saturday | t |
| tk_sunday | t |
| tk_attributes | 1 |
| tk_group | SERVER |
| tk_enable | t |
| tk_priority | 0 |

Built in Sensor & Tasks

| Sensor Name | Description |
|----------------------------|--|
| mon_command_history | Purges the command history table |
| mon_config | Saves any difference in the onconfig file |
| mon_config_startup | Save the onconfig file on every server startup |
| mon_profile | Save the server profile information |
| mon_vps | Collects the virtual processor timings |
| mon_checkpoint | Save information about checkpoints |
| mon_table_profile | Save table profile information including UDI counters |
| mon_table_names | Save the table names along with their create time |
| mon_users | Save profile information about each user |
| check_backup | Check to ensure backups have been done |

Admin Console

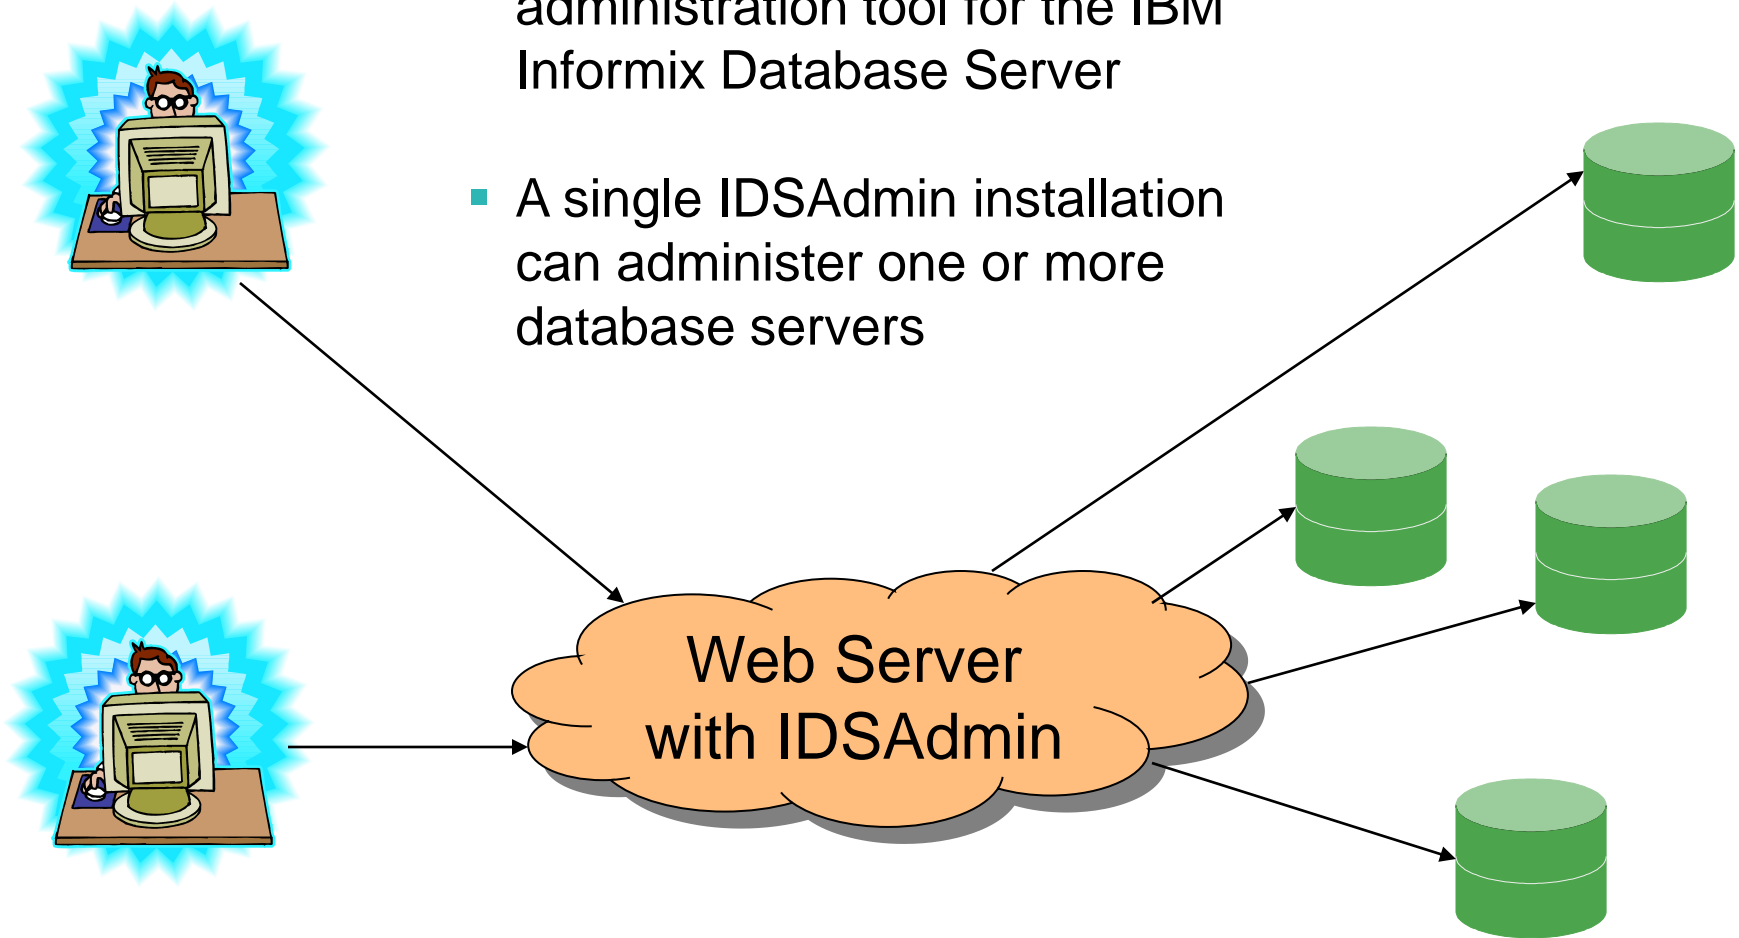


Configuring Scheduler

- **Scheduler is switched on by default**
- **Execute function task (“scheduler stop”)**
- **Execute function task (“scheduler start”)**
- **Execute function task(“reset sysadmin”, “dbspace1”);**
- **\$INFORMIXDIR/etc/sysadmin/stop – If present, scheduler will not start when the database server starts**

What is IDSAdmin?

- IDSAdmin is a web-based administration tool for the IBM Informix Database Server
- A single IDSAdmin installation can administer one or more database servers



Map View

Highlighting the server shows the information balloon

Home - Overview - Mozilla Firefox

File Edit View Go Bookmarks Tools Help

http://tugboat.beaverton.ibm.com/idsadmin/index.php

show treegraph.php

IDSAdmin

Connected: informix@jmiller_10wip
Host: olympia

Map Satellite Hybrid

Status: Online: CURRENT
Host: olympia
Informix Server: jmiller_10wip

Group Servers
jmiller_10adm
jmiller_10wip

Refresh

Task List
Run Times
Scheduler

Space
DBspaces
Chunks
Recovery Logs

Performance
SQL Trace
System Reports
User Reports

SQL ToolBox

RSS

Quick Info

Help

Logout

Info

Version: 10.50.
Boottime: 06-08 11:30
Time: 15:00:11
UpTime: 4 days 03:29:26
Sessions: 1
Max Sess: 10

Map data ©2006 TeleAtlas - Terms of Use

A Specific Instance of an SQL Statement

Query Tree

Query Statistics

SQL Profile

```

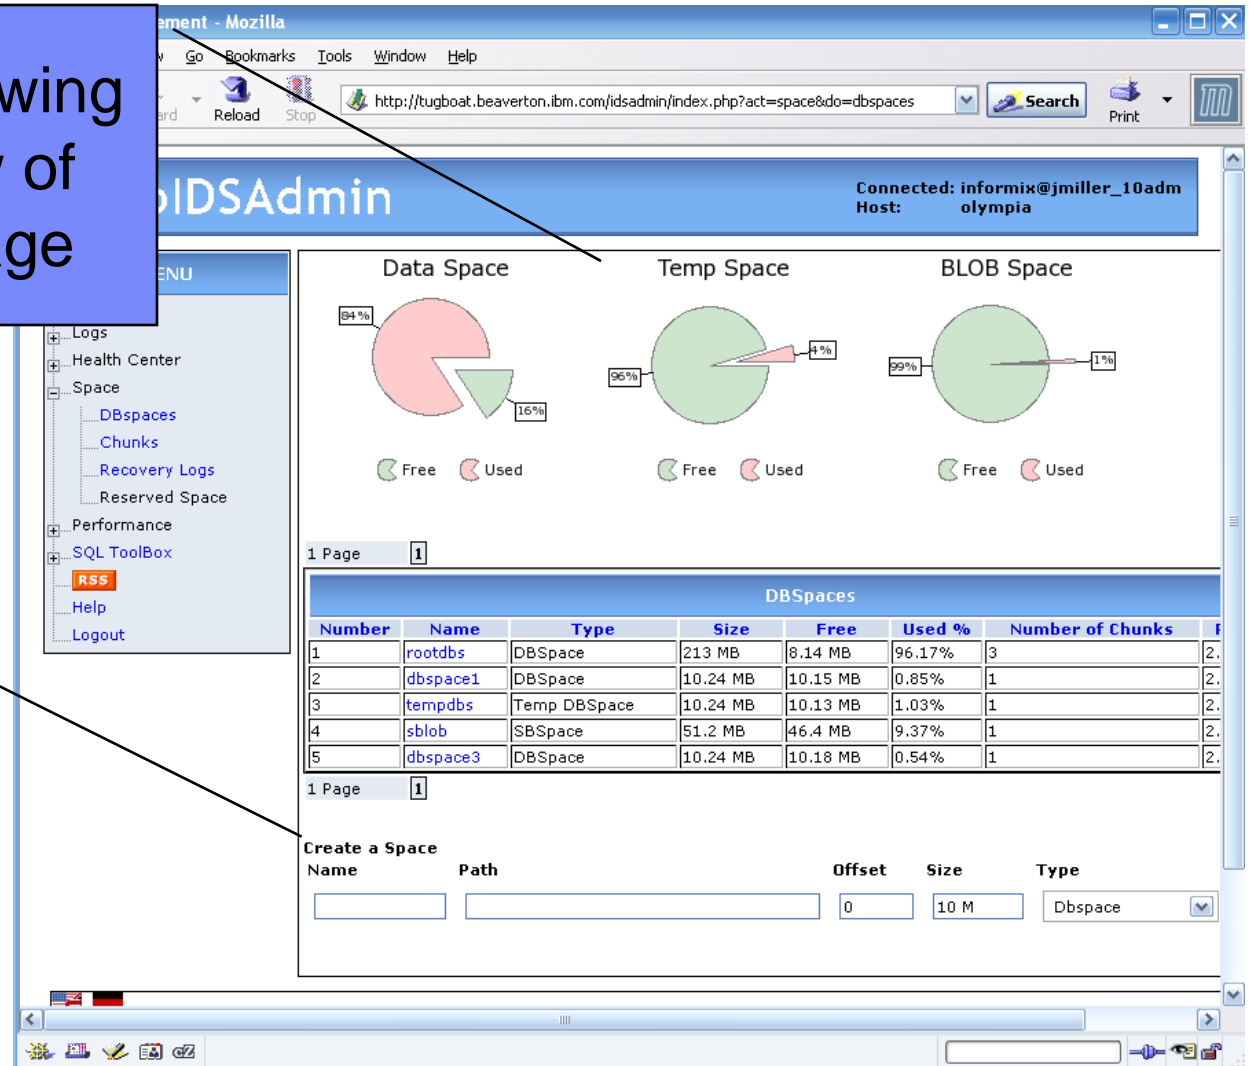
graph TD
    1[1.Group  
Cost 1  
Rows 1] --> 2[2.Hash Join  
Cost 36  
Rows 1]
    2 --> 3[3.Seq Scan  
Cost 8  
Rows 6]
    2 --> 4[4.Seq Scan  
Cost 8  
Rows 7]
  
```

| Session ID | User ID | Statement Type | Statement Completion Time | Response Time |
|------------|---------|----------------|---------------------------|---------------|
| 1666 | 900 | SELECT | 2006-06-20 17:02:23 | 0.06396695 |

| Statement Statistics | | | | | | | |
|----------------------|---------------------------|----------------------------|----------------------------|-------------------------|------------------|---------------------|---------------------|
| Page Reads | Buffer Reads | Reads Cache | Data Buffer Reads | Index Buffer Reads | Page Writes | Buffer Writes | Writes Cache |
| 0 | 2 | 100.00 % | 0 | 2 | 0 | 0 | 0.00 % |
| Lock Requests | # Lock Waits | Lock Wait Time (S) | Log Space | Disk Sorts | Memory Sorts | Number of Tables | Number of Iterators |
| 0 | 0 | 0 | 0.000 B | 0 | 0 | 2 | 4 |
| Total Executions | Total Executions Time (S) | Average Execution Time (S) | Maximum Execution Time (S) | Total Number of IO Wait | IO Wait Time (S) | Average Io Wait (S) | Average Rows/Second |
| 1 | 0.18879 | 0.18879 | 0.06396 | 0 | 0.00000 | 0.00000 | 15.63307 |
| Estimated Cost | Estimated Rows | Actual Rows | SQL Error | ISAM Error | Isolation Level | SQL Memory | |
| 36 | 1 | 1 | 0 | 0 | 2 | 99 KB | |

Overview of Disk Space

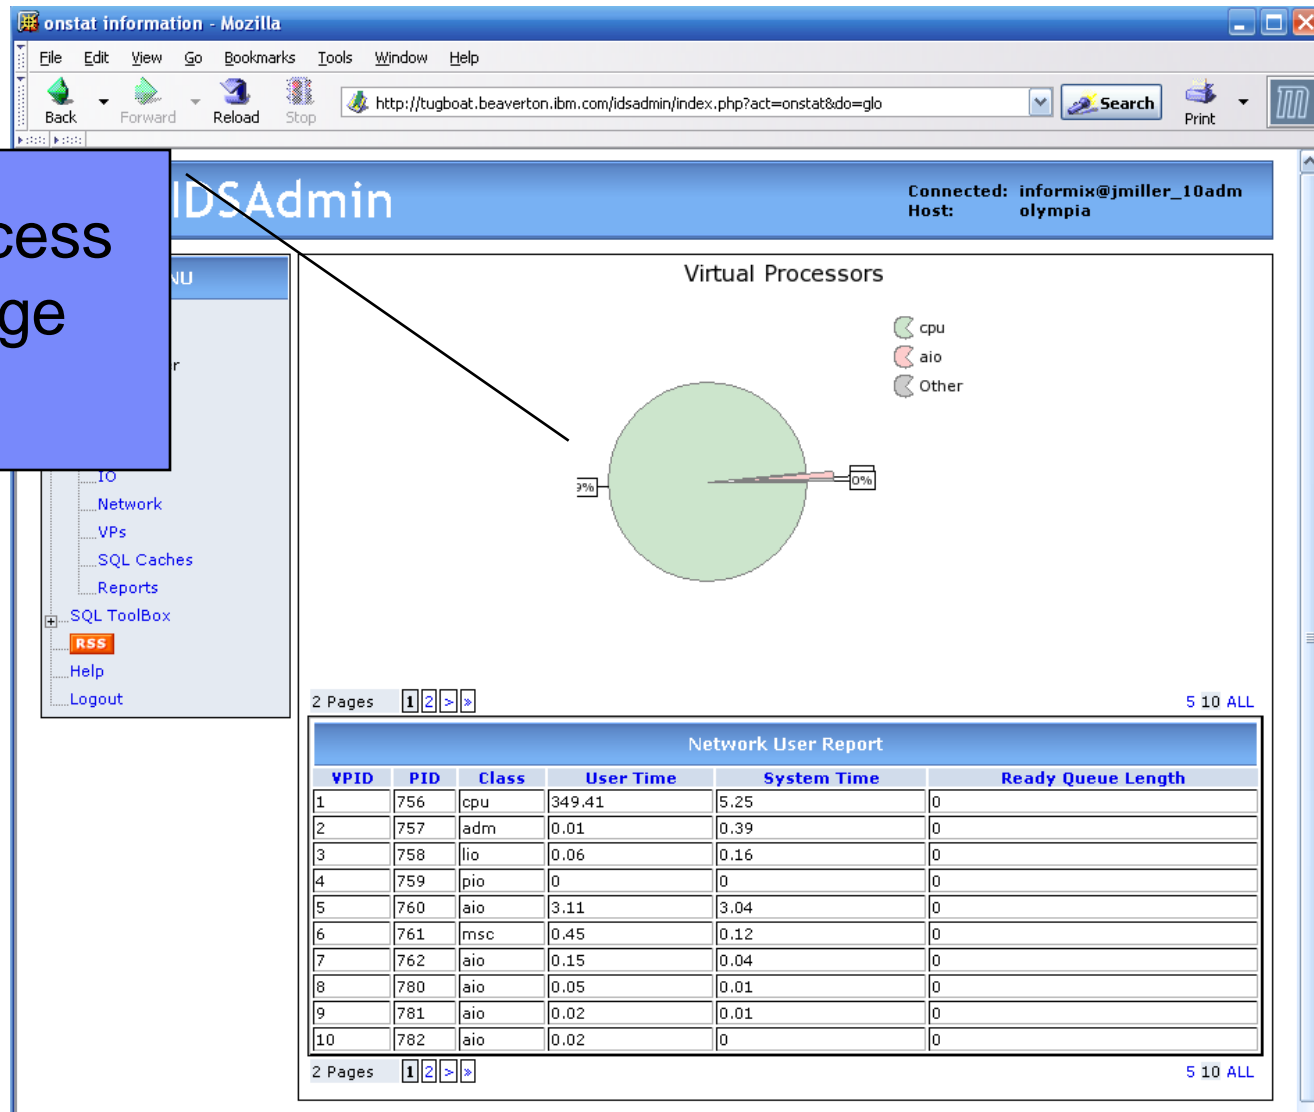
Graphs showing
quick view of
space usage



Ability to create
additional
spaces

Process Usage

Virtual Process CPU Usage Graph



Questions



Terminate Idle Users with the Database Admin System

- **The components we are going to utilize are:**

- Database Scheduler
- Alert System
- User Configurable Thresholds
- SQL Admin API

- **We are going to break the above problem into three separate parts.**

- Creating a tunable threshold for the idle time out
- Develop a stored procedure to terminate the idle users
- Schedule this procedure to run at regular intervals

INSERT INTO

```
ph_threshold(name,task_name,value,value_type,description)VALUES("IDLE  
TIMEOUT", "Idle Timeout","60","NUMERIC","Maximum amount of time in  
minutes for non-informix users to be idle.");
```

```
CREATE FUNCTION idle_timeout( task_id INT, task_seq INT)
RETURNING INTEGER
DEFINE time_allowed INTEGER;
DEFINE sys_hostname CHAR(16);
DEFINE sys_username CHAR(257);
DEFINE sys_sid    INTEGER;
DEFINE rc        INTEGER;
{*** Get the maximum amount of time to be idle ***}
SELECT value::integer    INTO time_allowed
FROM ph_threshold
WHERE name = "IDLE TIMEOUT";
{*** Find all users who are idle longer than the threshold ***}
FOREACH
SELECT admin("onmode","z",A.sid), A.username, A.sid, hostname
INTO rc, sys_username, sys_sid, sys_hostname
FROM sysmaster:sysrstcb A , sysmaster:systcblst B,sysmaster:sysscbllst C
WHERE A.tid = B.tid AND C.sid = A.sid AND lower(name) in ("sqlexec")
AND CURRENT - DBINFO("utc_to_datetime",last_run_time) > time_allowed UNITS MINUTE
AND lower(A.username) NOT IN( "informix", "root")
{*** If we successfully terminated a user log ***}
{*** the information into the alert table ***}
```

```
IF rc > 0 THEN
  INSERT INTO ph_alert(ID, alert_task_id,alert_task_seq, alert_type, alert_color, alert_state,
    alert_object_type, alert_object_name,alert_message,alert_action)
    VALUES (0,task_id, task_seq,"INFO", "GREEN", "ADDRESSED","USER","TIMEOUT",
      "User "||TRIM(sys_username)||"@ "||TRIM(sys_hostname)||
      " sid("||sys_sid||")"|| " terminated due to idle timeout.",NULL );
END IF
END FOREACH;
RETURN 0;
END FUNCTION;
```

```
/* *****
```

Create a task which will schedule the idle_timeout

SPL * to be run between 6 AM and 6 PM.

```
***** */
```

```
INSERT INTO
  ph_task(tk_name,tk_type,tk_group,tk_description,tk_execute,tk_start_time,tk_stop_time,tk_fr
    equency)
VALUES("Idle Timeout","TASK","USER","Remove all idle users from the
  system.","idle_timeout",DATETIME(06:00:00) HOUR TO SECOND,DATETIME(18:00:00)
  HOUR TO SECOND,INTERVAL ( 10 ) MINUTE TO MINUTE);
```